

Intro to Data Science - Lab 10

Copyright 2022, Jeffrey Stanton and Jeffrey Saltz Please do not post online.

Week 10 - Association Rules Mining

Enter your name here: Hendi Kushta

Please include nice comments.

Instructions:

Run the necessary code on your own instance of R-Studio.

Attribution statement: (choose only one and delete the rest)

1. I did this lab assignment by myself, with help from the book and the professor.

Association rules mining, also known as **market basket analysis**, is an **unsupervised data mining technique** that discovers patterns in the form of if-then rules. The technique is **** unsupervised **** in the sense that there is no prediction or classification happening. We are simply trying to find interesting **patterns**. In addition to working with baskets of objects, association rules mining is good at working with any kind of data that can be expressed as **lists of attributes**. For example, a trip to Washington DC might consist of the following attributes: train, July, morning departure, afternoon arrival, Union Station, first class, express.

In these exercises we will work with a built in data set called **groceries**. Make sure to library the **arules** and **arulesViz** packages before running the following:

```
data (Groceries) # Load data into memory
myGroc <- Groceries # Make a copy for safety
```

```
# install.packages("arules")
# install.packages("arulesViz")
```

```
library(arules)
```

```
## Loading required package: Matrix
```

```
##
```

```
## Attaching package: 'arules'
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      abbreviate, write
```

```
library(arulesViz)
```

```
data (Groceries) # Load data into memory
```

```
myGroc <- Groceries # Make a copy for safety
```

1. Examine the data structure that **summary()** reveals. This is called a **sparse matrix** and it efficiently stores a set of market baskets along with meta-data. Report using R comments about some of the item labels.

```
summary(myGroc)
```

```
## transactions as itemMatrix in sparse format with
## 9835 rows (elements/itemsets/transactions) and
## 169 columns (items) and a density of 0.02609146
```

```
##
## most frequent items:
##      whole milk other vegetables      rolls/buns      soda
##      2513      1903      1809      1715
##      yogurt      (Other)
##      1372      34055
##
## element (itemset/transaction) length distribution:
## sizes
##      1      2      3      4      5      6      7      8      9      10      11      12      13      14      15      16
## 2159 1643 1299 1005 855 645 545 438 350 246 182 117 78 77 55 46
##      17      18      19      20      21      22      23      24      26      27      28      29      32
##      29      14      14      9      11      4      6      1      1      1      1      3      1
##
##      Min. 1st Qu.  Median      Mean 3rd Qu.      Max.
##      1.000   2.000   3.000   4.409   6.000  32.000
##
## includes extended item information - examples:
##      labels level2      level1
## 1 frankfurter sausage meat and sausage
## 2      sausage sausage meat and sausage
## 3  liver loaf sausage meat and sausage
```

*# most frequent items are whole milk, vegetables, rolls/buns, sode,
yogurt. There is no basket without any item. minimum 1 item in a
basket and maximum 32. We can see that whole milk is bought more
frequently.*

2. Use the `itemFrequency(myGroc)` command to generate a list of item frequencies. Save that list in a new data object. Run `str()` on the data object and write a comment describing what it is. Run `sort()` on the data object and save the results. Run `head()` and `tail()` on the sorted object to show the most and least frequently occurring items. What s the most frequently purchased item?

```
newData <- itemFrequency(myGroc)
str(newData)

##      Named num [1:169] 0.05897 0.09395 0.00508 0.02603 0.02583 ...
##      - attr(*, "names")= chr [1:169] "frankfurter" "sausage" "liver loaf" "ham" ...
```

*# it shows the frequency of an item being purchased.
some of the items include "frankfurter" "sausage" "liver loaf" "ham" etc.*

```
newData <- sort(newData)
head(newData)

##      baby food  sound storage medium preservation products
##      0.0001016777      0.0001016777      0.0002033554
##      kitchen utensil      bags      frozen chicken
##      0.0004067107      0.0004067107      0.0006100661
```


```
tail(newData)

##      bottled water      yogurt      soda      rolls/buns
##      0.1105236      0.1395018      0.1743772      0.1839349
##      other vegetables      whole milk
##      0.1934926      0.2555160
```

the most frequent purchased item is whole milk.

3. Create a frequency plot with `itemFrequencyPlot(myGroc, topN=20)` and confirm that the plot shows the most frequently purchased item with the left-most bar. Write a comment describing the meaning of the Y-axis.

```
itemFrequencyPlot(myGroc, topN=20)
```



```
# it shows the relative frequency of occurrence of different items
# in the matrix for example the whole milk appears 2513 out of 9835 rows
# which means, 2513 / 9835 = 0.2555... or approximately 25.5% of the rows.
```

4. Create a cross table with `ct <- crossTable(myGroc, sort=TRUE)`. Examine the first few rows and columns of `ct` by using the square brackets subsetting technique. For example, the first two rows and first three columns would be `ct[1:2, 1:3]`. Write a comment describing one of values. Write a comment describing what is on the diagonal of the matrix.

```
ct <- crossTable(myGroc, sort=TRUE)
ct[1:2, 1:3]
```

```
##                whole milk other vegetables rolls/buns
## whole milk          2513             736         557
## other vegetables     736             1903         419
```

```
# the intersection of the same items with one another create a diagonal,
# which shows the real number of how many times the item is shown in the basket.
# the intersection of items with one another for example other vegetables and whole
# milk 736 means that we have bought vegetables and whole milk together this many times.
```

5. Run the following analysis:

```
rules1 <- apriori(myGroc,
  parameter=list(supp=0.0008, conf=0.55),
  control=list(verbose=F),
  appearance=list(default="lhs",rhs=("bottled beer")))
```

```
rules1 <- apriori(myGroc,
  parameter=list(supp=0.0008, conf=0.55),
  control=list(verbose=F),
```

```
appearance=list(default="lhs",rhs=("bottled beer"))
```

6. Examine the resulting rule set with `inspect()` and make sense of the results. There should be four rules in total.

```
inspect(rules1)
```

```
##      lhs                                rhs      support      confidence
## [1] {liquor, red/blush wine}      => {bottled beer} 0.0019318760 0.9047619
## [2] {soda, liquor}                => {bottled beer} 0.0012201322 0.5714286
## [3] {red/blush wine, napkins}     => {bottled beer} 0.0008134215 0.5714286
## [4] {soda, liquor, red/blush wine} => {bottled beer} 0.0008134215 1.0000000
##      coverage      lift      count
## [1] 0.0021352313 11.23527 19
## [2] 0.0021352313  7.09596 12
## [3] 0.0014234875  7.09596  8
## [4] 0.0008134215 12.41793  8
```

```
# There are huge chances that the items in the 4 lists, are bought together.
# confidence level, lift is very high in rule 1 and 4
# for rule 2 and 3 the confidence level is above 50%
```

7. Adjust the **support** parameter to a new value so that you get more rules. Anywhere between 10 and 30 rules would be fine. Examine the new rule set with `inspect()`. Does your interpretation of the situation still make sense?

```
rules2 <- apriori(myGroc,
  parameter=list(supp=0.0005, conf=0.55),
  control=list(verbose=F),
  appearance=list(default="lhs",rhs=("bottled beer")))
inspect(rules2)
```

```
##      lhs                                rhs      support confidence      coverage      lift count
## [1] {liquor (appetizer),
##      dishes}      => {bottled beer} 0.0006100661  0.8571429 0.0007117438 10.643939      6
## [2] {liquor,
##      red/blush wine}      => {bottled beer} 0.0019318760  0.9047619 0.0021352313 11.235269     19
## [3] {soda,
##      liquor}      => {bottled beer} 0.0012201322  0.5714286 0.0021352313  7.095960     12
## [4] {red/blush wine,
##      napkins}      => {bottled beer} 0.0008134215  0.5714286 0.0014234875  7.095960      8
## [5] {soda,
##      liquor,
##      red/blush wine}      => {bottled beer} 0.0008134215  1.0000000 0.0008134215 12.417929      8
## [6] {whole milk,
##      soups,
##      bottled water}      => {bottled beer} 0.0005083884  0.8333333 0.0006100661 10.348274      5
## [7] {yogurt,
##      pastry,
##      flower (seeds)}      => {bottled beer} 0.0005083884  0.8333333 0.0006100661 10.348274      5
## [8] {whole milk,
##      yogurt,
##      flower (seeds)}      => {bottled beer} 0.0005083884  0.7142857 0.0007117438  8.869949      5
## [9] {other vegetables,
##      salt,
##      margarine}      => {bottled beer} 0.0005083884  0.7142857 0.0007117438  8.869949      5
```

```
## [10] {soda,
##       red/blush wine,
##       napkins}          => {bottled beer} 0.0005083884 0.8333333 0.0006100661 10.348274 5
## [11] {citrus fruit,
##       oil,
##       bottled water}    => {bottled beer} 0.0005083884 0.5555556 0.0009150991 6.898850 5
## [12] {root vegetables,
##       herbs,
##       other vegetables,
##       bottled water}    => {bottled beer} 0.0006100661 0.6000000 0.0010167768 7.450758 6
## [13] {whole milk,
##       butter,
##       rolls/buns,
##       napkins}          => {bottled beer} 0.0005083884 0.5555556 0.0009150991 6.898850 5
## [14] {pork,
##       whole milk,
##       domestic eggs,
##       rolls/buns}       => {bottled beer} 0.0005083884 0.5555556 0.0009150991 6.898850 5
# now there are more rules, but still, the confidence level and lift are high.
```

8. Power User (not required): use **mtcars** to create a new data frame with **factors** (e.g., cyl attribute). Then create an mpg column with good or bad (good MPG is above 25). Convert the data frame to a transactions dataset and then predict rules for having bad MPG.

```
# data("mtcars")
# mtcars.t <- data.frame(mtcars)
# mtcars.t$mpg <- as.factor(mtcars.t$mpg)
# mtcars.t$cyl <- as.factor(mtcars.t$cyl)
# mtcars.t$disp <- as.factor(mtcars.t$disp)
# mtcars.t$hp <- as.factor(mtcars.t$hp)
# mtcars.t$drat <- as.factor(mtcars.t$drat)
# mtcars.t$wt <- as.factor(mtcars.t$wt)
# mtcars.t$qsec <- as.factor(mtcars.t$qsec)
# mtcars.t$vs <- as.factor(mtcars.t$vs)
# mtcars.t$am <- as.factor(mtcars.t$am)
# mtcars.t$gear <- as.factor(mtcars.t$gear)
# mtcars.t$carb <- as.factor(mtcars.t$carb)

# mtcars.t$goodOrBadMpg <- with(mtcars.t, ifelse(mpg > 25, 'Good', 'Bad'))
# str(mtcars.t)

# mtcars.t$mpg <- as.factor(mtcars.t$mpg)
# mtcars.t$goodOrBadMpg <- as.factor(mtcars.t$goodOrBadMpg)

# df_test <- mtcars.t[, c("mpg", "cyl", "goodOrBadMpg")]
# trans_df_test = as(df_test, "transactions")
# apriori(trans_df_test, parameter=list(supp=0.0004, conf=0.3))

# trans_mtcars.t = as(mtcars.t, "transactions")

# rulesMtgBadOrGood <- apriori(trans_mtcars.t, parameter=list(supp=0.0004, conf=0.3))

myCars <- mtcars
myCars <- data.frame(cyl = as.factor(mtcars$cyl),
```

```

gear = as.factor(mtcars$gear),
goodMpg = as.factor(mtcars$mpg > 25))

myCars$goodMpg

## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [13] FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE FALSE FALSE FALSE
## [25] FALSE TRUE TRUE TRUE FALSE FALSE FALSE FALSE
## Levels: FALSE TRUE

rules1 <- apriori(myCars, control = list(verbose=F),
                  parameter = list(supp=0.0001, conf=0.05),
                  appearance = list(default="lhs", rhs=("goodMpg=FALSE")))

inspect(rules1)

##      lhs                rhs      support confidence coverage lift
## [1] {}                  => {goodMpg=FALSE} 0.81250 0.8125000 1.00000 1.0000000
## [2] {gear=5}             => {goodMpg=FALSE} 0.09375 0.6000000 0.15625 0.7384615
## [3] {cyl=6}              => {goodMpg=FALSE} 0.21875 1.0000000 0.21875 1.2307692
## [4] {cyl=4}              => {goodMpg=FALSE} 0.15625 0.4545455 0.34375 0.5594406
## [5] {gear=4}             => {goodMpg=FALSE} 0.25000 0.6666667 0.37500 0.8205128
## [6] {cyl=8}              => {goodMpg=FALSE} 0.43750 1.0000000 0.43750 1.2307692
## [7] {gear=3}             => {goodMpg=FALSE} 0.46875 1.0000000 0.46875 1.2307692
## [8] {cyl=6, gear=5}      => {goodMpg=FALSE} 0.03125 1.0000000 0.03125 1.2307692
## [9] {cyl=8, gear=5}      => {goodMpg=FALSE} 0.06250 1.0000000 0.06250 1.2307692
## [10] {cyl=6, gear=4}     => {goodMpg=FALSE} 0.12500 1.0000000 0.12500 1.2307692
## [11] {cyl=6, gear=3}     => {goodMpg=FALSE} 0.06250 1.0000000 0.06250 1.2307692
## [12] {cyl=4, gear=4}     => {goodMpg=FALSE} 0.12500 0.5000000 0.25000 0.6153846
## [13] {cyl=4, gear=3}     => {goodMpg=FALSE} 0.03125 1.0000000 0.03125 1.2307692
## [14] {cyl=8, gear=3}     => {goodMpg=FALSE} 0.37500 1.0000000 0.37500 1.2307692
##      count
## [1] 26
## [2] 3
## [3] 7
## [4] 5
## [5] 8
## [6] 14
## [7] 15
## [8] 1
## [9] 2
## [10] 4
## [11] 2
## [12] 4
## [13] 1
## [14] 12

```